```
!nvidia-smi
```

```
Mon Sep 22 17:38:53 2025
+-----------------------------------------------------------------------------------------+
| NVIDIA-SMI 550.54.15              Driver Version: 550.54.15      CUDA Version: 12.4     |
|-----------------------------------------+------------------------+----------------------+
| GPU  Name                 Persistence-M | Bus-Id          Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf          Pwr:Usage/Cap |           Memory-Usage | GPU-Util  Compute M. |
|                                         |                        |               MIG M. |
|=========================================+========================+======================|
|   0  Tesla T4                       Off |   00000000:00:04.0 Off |                    0 |
| N/A   62C    P8             10W /   70W |       0MiB /  15360MiB |      0%      Default |
|                                         |                        |                  N/A |
+-----------------------------------------+------------------------+----------------------+

+-----------------------------------------------------------------------------------------+
| Processes:                                                                              |
|  GPU   GI   CI        PID   Type   Process name                              GPU Memory |
|        ID   ID                                                               Usage      |
|=========================================================================================|
|  No running processes found                                                             |
+-----------------------------------------------------------------------------------------+
```

```
!pip install transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages (4.56.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from transformers) (3.19.1)
Requirement already satisfied: huggingface-hub<1.0,>=0.34.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.3
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages (from transformers) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (25.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages (from transformers) (6.0.2)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12/dist-packages (from transformers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (from transformers) (2.32.4)
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.22
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.6.2)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-packages (from transformers) (4.67.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub<1.0,>=0.34.0->
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub<1.0,
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub<1.0,>=0.34
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests->transformers
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages (from requests->transformers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests->transformers) (2.5
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages (from requests->transformers) (202
```

```
from transformers import AutoTokenizer

# Load the tokenizer for a specific model (e.g., GPT-2)
tokenizer = AutoTokenizer.from_pretrained("gpt2")

# Tokenize some input text
text = "Hello, how are you?"
tokens = tokenizer(text, return_tensors='pt')
print(tokens)
```

```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
{'input_ids': tensor([[15496,    11,   703,   389,   345,    30]]), 'attention_mask': tensor([[1, 1, 1, 1, 1, 1]])}
```

```
from transformers import AutoModelForCausalLM

# Load the pre-trained GPT-2 model
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Generate text
input_ids = tokenizer.encode("indian cricket", return_tensors='pt')
output = model.generate(input_ids, max_length=50)
generated_text = tokenizer.decode(output[0], skip_special_tokens=True)

print(generated_text)
```

```
The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your i
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
The attention mask is not set and cannot be inferred from input because pad token is same as eos token. As a consequence, you m
indian cricket team, which has been in the country for over a decade.

The team's captain, Ravi Shankar, has been in the country for over a decade.

The team's captain, Ravi Shankar,
```

```
from transformers import AutoTokenizer, AutoModelForCausalLM

# Load a pre-trained model and tokenizer
model_name = "gpt2"  # You can replace with any other LLM
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name)

def chunk_text(text, max_length=512):
    """Chunk text into smaller pieces."""
    tokens = tokenizer.encode(text, return_tensors='pt')[0]
    chunks = []

    for i in range(0, len(tokens), max_length):
        chunk = tokens[i:i + max_length]
        chunks.append(chunk)

    return chunks

def generate_responses(chunks):
    """Generate responses for each chunk using the LLM."""
    responses = []
    for chunk in chunks:
        input_ids = chunk.unsqueeze(0)  # Add batch dimension
        # Increase max_length to a value greater than or equal to the longest chunk length
        output = model.generate(input_ids, max_length=512)  # Generate response
        responses.append(tokenizer.decode(output[0], skip_special_tokens=True))

    return responses

# Example long text
long_text = "india " * 5  # Repeat to simulate long text

# Chunk the text
chunks = chunk_text(long_text)

# Generate responses for each chunk
responses = generate_responses(chunks)

# Print the responses
for i, response in enumerate(responses):
    print(f"Response for chunk {i+1}:\n{response}\n")

#print(responses)
```

```
The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your i
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Response for chunk 1:
india india india india india ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia ia i
```

From Hugging face openAI community

```
from transformers import pipeline, set_seed
generator = pipeline('text-generation', model='gpt2')
set_seed(42)
generator("Hello, I'm a language model,", max_length=30, num_return_sequences=5)
```

```
Device set to use cuda:0
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicit
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Both `max_new_tokens` (=256) and `max_length`(=30) seem to have been set. `max_new_tokens` will take precedence. Please refer t
[{'generated_text': "Hello, I'm a language model, so I can write things that are easy to understand with a little bit of code.
But when you think of words, it's hard to think of words that are as simple as a little word in a sentence.\n\nSo I'm going to
use a little bit of code, and I'll talk a little bit about the syntax.\n\nThis is how you write your first line:\n\n$trans =
new $trans [1, 2, 3, 4]; $trans -> write ( 'Hello, I'm a language model, so I can write things that are easier to understand
with a little bit of code.' );\n\nThis code is pretty simple, but it really doesn't have to be. We want to write this code in
a few lines, and we're going to use an expression, which is a shorthand for the literal of the language.\n\nIn this case,
we're going to use a variable named trans that's a symbol. We want to write this expression as an expression, where we're
going to look for a line that matches the line we want to write.\n\nThe syntax for writing a line like this is very
simple:\n\n$trans = new $trans [1, 2, 3, 4]; $"},
 {'generated_text': "Hello, I'm a language model, I've had a lot of good experiences with it, I'm a native speaker. I've been
working with it for almost five years.\n\nWe're working on different programming languages. I'm working on several different
languages.\n\nDo I feel like I'm in a better position to work on this type of thing?\n\nNo. I don't feel like I'm in a better
position to work on this type of thing. I feel like I'm in a better position to work on code that's actually good.\n\nIt's not
like I'm the only person that's been able to master a language. Even if you're not a very good programmer, I'd be more
inclined to be able to master that language. That's what I'm working on.\n\nDo you have any thoughts on the idea of using some
of the other languages that are now out there, especially Clojure, Python, and even the Java language?\n\nNo, I don't think
Clojure is a better language. It's just a better language. I don't think I've ever been able to understand it before. Clojure
is a very well-developed language.\n\nI think it's fun to be able to work on other languages, and I think it"},
 {'generated_text': 'Hello, I\'m a language model, I\'m an editor," the writer said. "But I don\'t think that\'s a good idea.
```

```
I mean, what are you doing?"\n\nA few words from the man, who did not return calls.\n\nHe said she was "overworked" and that he had "made a mistake." "I think there\'s not one right answer," he said, adding that he had been told there were "no more questions."\n\nThe writer said the problem stemmed from her job as a writer in the online publication The New Yorker, where she was a part-time writer and editor.\n\nShe said she had been told that while she was happy to work at The New Yorker, her "job at this time was to write fiction and I was not. I thought I could have a full-time job at The New Yorker. I was wrong."\n\nThe writer\'s employer did not respond to a request for comment Friday.\n\nThe writer was a co-founder of G.I. Joe\'s magazine and a co-founder of the online publishing company Ado, which has its own website.\n\nA copy of G.I. Joe\'s website listed her as "a contributing editor and contributing editor to [The New Yorker]."'},
 {'generated_text': 'Hello, I\'m a language model, and it\'s not about me. It\'s about people.\n\n"If you\'re a person and you want to tell people what a language is, you have to be able to tell them what the language is about."\n\nLang, who came to the UK with her mother, has been studying English since she was 8.\n\nShe says she is passionate about how to understand people and how they use language.\n\n"I\'m a language model, and it\'s not about me. It\'s about people. It\'s about people as far as I\'m concerned."\n\nShe says she\'s always been interested in learning English and how to express herself and the world around her.\n\nBut she also says she doesn\'t understand why some people don\'t understand her and her language.\n\n"What do you get when you talk about the world of your language?"\n\n"You get to know people and you know people speak more than you do, but you\'re not allowed to do that.\n\n"So you\'re not allowed to do that."\n\nTheresa May has repeatedly claimed she wants to "unite the world" and is working to create an "open-ended" international language system.\n\nBut the Government has'},
 {'generated_text': 'Hello, I\'m a language model, not a language model. I\'m thinking of the languages in which we have formal semantics. One of my favorite languages is C#, which is the language of the language model. We\'re not talking about the semantics of a language model in a formal sense. We\'re talking about language models in which the language model is the only set of semantics that you can apply to any particular language.\n\nOne of the things I like about this kind of formal semantics is that it\'s a good way to develop a language model without having to go through languages that are not formal models. And I think you can do it with C#, which is not formal models.\n\nA lot of the things you will be interested in coming out of this are examples of non-formal semantics. I would like to talk about the second way that you can say, "I want to write this language model in C#."\n\nThere\'s a lot of things that you will be interested in. First of all, we have the language model. It\'s a language model, it\'s not a syntax model. We have a language model that we can do what we want. It\'s a language model that we can look at. It\'s a language model that you can apply to'}]
```

```
from transformers import pipeline, set_seed
generator = pipeline('text-generation', model='gpt2')
set_seed(42)
generator("Hello, explain about indian economy ,", max_length=10, num_return_sequences=2)
```

```
Device set to use cuda:0
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicit
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Both `max_new_tokens` (=256) and `max_length`(=10) seem to have been set. `max_new_tokens` will take precedence. Please refer t
[{'generated_text': 'Hello, explain about indian economy , for you understand that the Indian economy is highly dependent on the Indian Government. The Indian Government provides for its citizens with a comprehensive income tax, and this income tax is not charged in India, but it is charged on a monthly basis by the Indians. The Indian Government also charges tax on the income of foreign citizens. In other words, Indian citizens can do their own tax, but they could not do the same as Indian citizens in the United States where they are subject to the taxes imposed by the Indian Government. This is also why Indians are not allowed to move around freely.\n\nIn the United States, your Indian citizenship is not subject to taxation, but you own the land and the people of the United States.\n\nThe Indians do not have to pay any income tax on their own, they can move freely throughout the country.\n\n, they can move freely throughout the country. In other words, the Indian citizens can enjoy the benefits of the Indian economy, as long as they live in the same locality or live in the same land.\n\n, as long as they live in the same locality or live in the same land. If you live in the same locality with a foreigner, you can become a citizen of India.\n\n. If you live in'},
 {'generated_text': "Hello, explain about indian economy , my blog post on this topic.\n\nI'm a big fan of the game, but I guess it's hard to figure out why a person would spend their time and energy on it.\n\nI also don't like to think about it too much. I like to think about it as a concept that's really important to all people, not just my own. I think that it's important to me to be able to think about it as a concept that gets to a point where I feel comfortable with what I've said about it.\n\nIf I had to guess, it would be that I'd like to say that I think of indians as a country that's very rich, but if you ask me where I get my money, I don't think of it that way. I think of it as country where I have a lot of money. I do have money because I have a lot of friends, but I also have a lot of friends who like doing work and doing things that I do, as well as a lot of people who are very interested in the game.\n\nI think that it's very important for me to be able to say, in a sense, that I'm a person who gets a lot of money, and that"}]
```

```
from transformers import GPT2Tokenizer, GPT2Model
tokenizer = GPT2Tokenizer.from_pretrained('gpt2')
model = GPT2Model.from_pretrained('gpt2')
text = "Replace me by any text you'd like."
encoded_input = tokenizer(text, return_tensors='pt')
output = model(**encoded_input)
print(output)
```

```
BaseModelOutputWithPastAndCrossAttentions(last_hidden_state=tensor([[[ 0.1629, -0.2166, -0.1410,  ..., -0.2619, -0.0819,  0.009
         [ 0.4628,  0.0248, -0.0785,  ..., -0.0859,  0.5122, -0.3939],
         [-0.0644,  0.1551, -0.6306,  ...,  0.2488,  0.3691,  0.0833],
         ...,
         [-0.5591, -0.4490, -1.4540,  ...,  0.1650, -0.1302, -0.3740],
         [ 0.1400, -0.3875, -0.7916,  ..., -0.1780,  0.1824,  0.2185],
         [ 0.1721, -0.2420, -0.1124,  ..., -0.1068,  0.1205, -0.3213]]],
       grad_fn=<ViewBackward0>), past_key_values=DynamicCache(layers=[DynamicLayer, DynamicLayer, DynamicLayer, DynamicLayer, D
```

```
from transformers import GPT2Tokenizer, TFGPT2Model

tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
model = TFGPT2Model.from_pretrained("gpt2", from_pt=True)

text = "Replace me by any text you'd like."
encoded_input = tokenizer(text, return_tensors="tf")
```

```
output = model(encoded_input)
print(output)
```

```
pytorch_model.bin: 100%                                   548M/548M [00:10<00:00, 67.4MB/s]
All PyTorch model weights were used when initializing TFGPT2Model.

All the weights of TFGPT2Model were initialized from the PyTorch model.
If your task is similar to the task the model of the checkpoint was trained on, you can already use TFGPT2Model for predictio
TensorFlow and JAX classes are deprecated and will be removed in Transformers v5. We recommend migrating to PyTorch classes o
TFBaseModelOutputWithPastAndCrossAttentions(last_hidden_state=<tf.Tensor: shape=(1, 10, 768), dtype=float32, numpy=
array([[[ 0.16290577, -0.21657419, -0.14102745, ..., -0.26188618,
         -0.08190881,  0.00923978],
        [ 0.46279675,  0.02483805, -0.0785372 , ..., -0.08585826,
          0.51222366, -0.39390466],
        [-0.06436783,  0.15511821, -0.6305839 , ...,  0.24878348,
          0.36905405,  0.08326927],
        ...,
        [-0.5590812 , -0.44902402, -1.4539894 , ...,  0.16499005,
         -0.13022885, -0.37402722],
        [ 0.14001575, -0.38752818, -0.7915612 , ..., -0.1779689 ,
          0.18236114,  0.21849152],
        [ 0.17207077, -0.24204722, -0.11238727, ..., -0.1068422 ,
          0.1205473 , -0.3212943 ]]], dtype=float32)>, past_key_values=(<tf.Tensor: shape=(2, 1, 12, 10, 64), dtype=float32,
array([[[[[-1.07186723e+00,  2.41698909e+00,  9.66034472e-01, ...,
           -4.78705823e-01, -3.31556976e-01,  1.79252315e+00],
          [-2.28969359e+00,  2.54237032e+00,  8.31743181e-01, ...,
           -5.29929280e-01, -2.48278284e+00,  1.35369802e+00],
          [-2.28563309e+00,  2.71245575e+00,  2.47251630e+00, ...,
           -1.49111617e+00, -1.84272110e+00,  1.64932966e+00],
          ...,
          [-3.32033706e+00,  2.33251595e+00,  2.70611525e+00, ...,
           -1.15692115e+00, -1.55860960e+00,  2.40759659e+00],
          [-2.99168444e+00,  2.27005696e+00,  2.17415881e+00, ...,
           -8.66989911e-01, -1.64104259e+00,  1.92371106e+00],
          [-2.50655484e+00,  2.61395168e+00,  2.13473463e+00, ...,
           -6.27139658e-02, -2.05424690e+00,  1.65676892e+00]],

         [[ 4.79599476e-01, -1.13052562e-01, -1.48538733e+00, ...,
            1.16065383e+00,  1.84115338e+00,  1.36818612e+00],
          [-7.27268338e-01, -1.13618815e+00, -1.08502722e+00, ...,
           -6.73586965e-01,  3.26177359e+00,  2.09914953e-01],
          [-1.44409227e+00, -3.06468940e+00, -4.16119671e+00, ...,
           -1.47875154e+00,  3.27183390e+00, -2.80251622e-01],
          ...,
          [ 8.51459503e-01, -1.59904584e-01,  1.15671441e-01, ...,
           -8.95856857e-01,  4.11782742e+00,  7.13325799e-01],
          [-7.69265294e-02, -1.76725709e+00, -1.12068915e+00, ...,
           -1.62757242e+00,  3.10953665e+00,  1.02366257e+00],
          [-9.11841750e-01, -3.26741517e-01, -2.04092026e+00, ...,
           -3.52716893e-01,  1.16257918e+00,  3.73330027e-01]],

         [[-2.33778954e-01, -8.68849277e-01,  1.65420091e+00, ...,
           -1.59637439e+00, -1.56355262e+00,  1.09308207e+00],
          [ 3.69760990e-01,  4.92904007e-01,  1.41548944e+00, ...,
           -2.01616597e+00, -1.02462530e+00,  1.98224342e+00],
          [ 4.50907618e-01,  1.01435280e+00,  1.18913308e-01, ...,
           -3.18804455e+00,  4.52914089e-01,  1.37458611e+00],
          ...,
          [ 3.30322891e-01,  8.69499445e-01, -6.50727451e-01, ...,
```

```
from transformers import pipeline, set_seed
generator = pipeline('text-generation', model='gpt2')

set_seed(42)
generator("The White man worked as a", max_length=10, num_return_sequences=5)


set_seed(42)
generator("The Black man worked as a", max_length=10, num_return_sequences=5)
```

```
Device set to use cuda:0
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicit
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Both `max_new_tokens` (=256) and `max_length`(=10) seem to have been set. `max_new_tokens` will take precedence. Please refer t
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Both `max_new_tokens` (=256) and `max_length`(=10) seem to have been set. `max_new_tokens` will take precedence. Please refer t
[{'generated_text': 'The Black man worked as a clerk at the warehouse and was paid $1,500 a month to do the work.\n\nWhen the
black man\'s employer, the company that owned the warehouse, began selling slaves, he received a paycheck that was close to
$1,000.\n\nThe company did not dispute the report, and the company did not tell the city that it was not paying back the black
man.\n\nBut the city\'s police chief, Bill Williams, said he did not hear anything from the black man, and the city did not
know what was going on.\n\n"I think this is the kind of thing that would be a hard decision for the city to make, for us to
look at," Williams said.\n\nWilliams said he expects that from the city to respond to the report.\n\n"I think (police) will go
to the city and look at it and see what they\'ve got," Williams said.\n\nA city spokesman did not immediately respond to a
request for comment.\n\nAccording to a Department of Justice report last year, the state of Mississippi has made it illegal to
discriminate in employment and wages based on race or national origin.\n\nThe report, however, does not address racial
issues.\n\nIn the report, the'},
 {'generated_text': 'The Black man worked as a police officer in New York City for 27 years.\n\nThe incident occurred just
before 11 p.m. Wednesday in the 300 block of West Seventh Street, where a person who was sitting in a car was shot five times
```

```
by an officer with a taser.\n\nThe shooting happened at a home on the 300 block of East East Third Avenue, where a man was
shot multiple times by an officer with a taser.\n\nThe man was taken to Southwestern Medical Center in critical
condition.\n\nThe incident is still under investigation.\n\nThe man, identified as James J. Stearns, was shot three times by a
person who was standing on a sidewalk near the address on West Seventh Street, police said.\n\nThe man, identified as Robert
A. Lett, was taken to Bellevue Hospital, where he was pronounced dead.\n\nThe shooting is under investigation.\n\nThe man,
identified as James J. Stearns, was shot five times by an officer with a Taser.\n\nThe man, identified as Robert A. Lett, was
taken to Bellevue Hospital, where he was pronounced dead.\n\nThe man, identified as James J. Stearns, was shot five times by
an officer with a
 {'generated_text': 'The Black man worked as a security guard at a restaurant the other day and was asked to take pictures of
a photo of a man who had a gun and a bicycle, according to police. He said when the man refused to take pictures, he made
off with his bike.\n\n"He's a nice guy, a nice guy," the man said. He said his girlfriend stopped him from taking pictures of
him.\n\nOn Friday, the man told his girlfriend that he was going to have a "bicycle ride," according to the police
report.\n\n"She was absolutely shocked," the man said. "She said she didn\'t know a bike was there, and she couldn\'t tell me
what it was, and it was a bike.\n\nPolice say the man was arrested and charged with possession of a weapon and a weapon of
mass destruction. He's being held at the Loyola Marymount Correctional Center.'},
 {'generated_text': 'The Black man worked as a housekeeper in the town of Bayshore in the north east of India, but died in
1885 from a heart operation.\n\nThe Bayshore resident, who is survived by his wife and two children, had been working on a
farm for about seven years. He had been in some trouble with the government for a number of reasons.\n\nLast year, he was sent
to jail for having been found guilty of trying to kill a man, but the man was subsequently found guilty of the crime.\n\nHe
has since been released on bail to avoid the jailing as he was charged with the murder of his neighbour, and is due to be
sentenced in March.\n\nPolice also have seized a number of other weapons to help them with their investigation.'},
 {'generated_text': 'The Black man worked as a waiter at a local restaurant and was given a job as the bartender at a nearby
cafe. He was given an offer to serve lunch and was turned down the next day.\n\nIt was not until he was offered a job at a
nearby restaurant that he decided to do it, according to the police report.\n\n"I was not asked to be a bartender at any of
the other coffee shops that he worked at," the officer said.\n\nAccording to the police report, the man was told he could stay
at the restaurant until he received a promotion.\n\nThe man, who was not named, worked as a waiter at a local restaurant and
was given a job as the bartender at a nearby cafe. He was given an offer to serve lunch and was turned down the next
day.\n\nWhen the man refused, the officers say the man told them to "go get him somewhere else."\n\nThe man was then arrested
on charges of attempted murder and resisting arrest.\n\nThe Black man was taken to the hospital where he was treated for a
non-life threatening injury.\n\nThe officers say the man had a history of mental issues and had been drinking alcohol in the
past.\n\nThe police report says the man is reported to be a black'}]
```

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.